Tracking Performance Drift in Large Language Models Across Successive Version Updates

Abstract

The integration of Large Language Models (LLMs) into critical enterprise, scientific, and consumer infrastructures relies upon an implicit assumption of temporal stability—specifically, that sequential version updates will strictly improve or maintain a model's foundational capabilities. A growing corpus of empirical evidence fundamentally challenges this assumption, revealing that commercial and open-weights LLMs exhibit significant, unpredictable variations in output distribution, reasoning capacity, and instruction adherence over time. This phenomenon, formalized as performance drift, introduces severe vulnerabilities regarding the reproducibility of AI-assisted workflows and the referential security of compliance mechanisms. This comprehensive research report investigates the architectural mechanics and empirical manifestations of LLM performance drift. By synthesizing literature on catastrophic forgetting during sequential instruction tuning, the alignment tax, and the scaling laws of reward model overoptimization, this paper delineates how continuous fine-tuning degrades generalized reasoning capabilities. Furthermore, the analysis evaluates state-of-the-art diagnostic and benchmarking frameworks that have emerged to quantify this instability, including dynamic contamination-resistant platforms (LiveBench), symbolic template evaluation for mathematical reasoning (GSM-Symbolic), and automated assessment paradigms (LLM-as-a-Judge). Finally, the report explores algorithmic interventions—ranging from Entropy-Adaptive Fine-Tuning and survival analysis in multimodal trajectories to neuro-symbolic execution engines—that aim to mitigate destructive interference and stabilize the continuous deployment of artificial intelligence systems.

Keywords

Performance Drift, Large Language Models, Catastrophic Forgetting, Alignment Tax, Reward Model Overoptimization, Continuous Evaluation, LLM-as-a-Judge, LiveBench, GSM-Symbolic, Referential Security.

Introduction

The operational paradigm of Large Language Models (LLMs) has transitioned from the deployment of static, monolithic software artifacts to continuously updated, dynamically served application programming interfaces (APIs). While continuous integration enables developers to rapidly issue safety patches, expand domain-specific knowledge bases, and refine human-alignment protocols, it concurrently introduces a critical systemic vulnerability known as performance drift. Performance drift refers to the longitudinal fluctuation in a neural network's output distribution, factual accuracy, and behavioral alignment when evaluated under identical prompt conditions across successive temporal versions[1].

The instability of these models challenges the core tenets of software reliability. In deterministic computing, iterative updates generally append new functionalities without compromising the structural integrity of existing operations. Neural networks, however, exhibit non-linear and highly coupled responses to continual post-training adjustments. A modification intended to enhance an LLM's refusal rate for adversarial prompts may inadvertently degrade its capacity for complex logical reasoning—an artifact of the "alignment tax"[2]. Similarly, an update designed to make a model's conversational tone more helpful can erode its strict adherence to syntactic formatting, rendering it incapable of generating clean, executable code[4]. The assumption of time invariance—that a model's average performance under fixed conditions remains stable over time—has been empirically disproven. Longitudinal studies, such as those evaluating a fixed mathematical physics task queried against a frontier model ten times every three hours over three months, demonstrate substantial variations in zero-shot capability[6].

The opacity of commercial frontier models severely exacerbates the challenge of tracking performance drift. AI laboratories routinely update model weights, adjust reinforcement learning parameters, and dynamically route queries through unannounced Mixture of Experts (MoE) architectures without detailing the specific modifications or their anticipated side effects[1]. For regulated workflows spanning healthcare diagnostics, quantitative finance, and public policy administration, this silent drift transforms the boundary of compliance into a moving target[8]. Consequently, the academic and engineering communities have been compelled to develop adversarial, dynamically updating evaluation frameworks to track and mitigate the degradation of model capabilities over time. This report provides an exhaustive, peer-reviewed examination of the scholarly literature concerning LLM performance drift, analyzing empirical evidence, underlying mathematical causes, evolving evaluation methodologies, and future directions for establishing referential stability within artificial intelligence ecosystems.

Background and Literature Review

To comprehensively understand performance drift, it is necessary to examine the underlying mechanisms of neural network training that induce instability. The literature identifies three primary vectors for performance degradation during the post-training lifecycle of an LLM: catastrophic forgetting during supervised fine-tuning, the alignment tax, and reward model overoptimization.

Catastrophic Forgetting During Continual Fine-Tuning

Catastrophic forgetting (CF) represents a foundational challenge in connectionist architectures. Initially documented in early neural network research, CF occurs when the sequential learning of new tasks fundamentally overwrites the parameter weights associated with previously acquired knowledge, resulting in severe performance degradation on earlier objectives[2]. In the context of modern LLMs, catastrophic forgetting manifests prominently during Supervised Fine-Tuning (SFT) and domain-specific adaptation[2].

An extensive empirical investigation into this phenomenon is detailed in *An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning* by Luo et al., which systematically evaluated the forgetting phenomenon during the continual instruction tuning of open-source LLMs ranging from 1 billion to 7 billion parameters[3]. The researchers assessed knowledge retention across distinct vectors, including domain knowledge, reasoning capabilities, and reading comprehension, utilizing sequential fine-tuning on diverse generative tasks (e.g., text simplification, empathetic dialogue generation, inquisitive question generation)[9]. The findings established a counterintuitive scaling law regarding catastrophic forgetting: as the parameter scale of the model increases within the evaluated range, the severity of forgetting intensifies[9]. The authors hypothesize that larger models achieve significantly higher initial performance on generalized tasks, which provides a much steeper gradient for performance degradation when the model's weights are aggressively updated to accommodate narrow, task-specific instruction sets[9]. This observation was independently corroborated in evaluations of the Qwen-2.5-Instruct architecture scaling from 3B to 14B parameters[3].

Architectural variations also dictate the severity of CF. When comparing the decoder-only BLOOMZ model with the encoder-decoder mT0 model, BLOOMZ exhibited substantially less forgetting and superior knowledge retention when subjected to identical fine-tuning datasets[9]. Furthermore, the study indicated that general instruction tuning acts as a protective buffer; base models that had undergone broad, diverse instruction tuning (such as the transition from LLaMA to Alpaca) prior to sequential specialized fine-tuning experienced significantly less knowledge degradation than those adapted directly from pre-trained weights[9]. Across knowledge domains, reading comprehension was identified as the most susceptible to severe forgetting, whereas logical reasoning capabilities demonstrated slightly higher resilience[9].

To combat CF, researchers have proposed various architectural interventions. Traditional regularization techniques like Elastic Weight Consolidation (EWC) attempt to penalize changes to parameters deemed critical for past tasks. Recent advancements, however, have introduced the Range Matrix, which improves memory retention and reduces CF relative to EWC by utilizing static fully connected architectures more efficiently[13]. Parameter-efficient fine-tuning (PEFT) methods, such as Low-Rank Adaptation (LoRA), adapt LLMs cheaply but suffer from CF when low-rank updates are continuously stacked on frozen weights, causing new tasks to overwrite predecessors[13]. The ReCoLoRA algorithm addresses this by recursively re-decomposing the effective weight into a frozen residual, a slowly updated principal component, and a fresh adapter prior to each new task, thereby allowing every task to build upon a consolidated model[13]. Similarly, the SA-SFT framework introduces a lightweight self-augmentation routine where an LLM generates self-dialogues prior to fine-tuning, mixing self-authored data with task data to mitigate CF without modifying the optimization schedule[13].

Reward Model Overoptimization and the Alignment Tax

The second major driver of performance drift originates within the Reinforcement Learning with Human Feedback (RLHF) pipeline. During RLHF, an LLM generates responses that are optimized against a parameterized reward model acting as a proxy for human preference[14]. In their seminal research, *Scaling Laws for Reward Model Overoptimization*, Gao, Schulman, and Hilton demonstrate that the relationship between optimization pressure and actual model performance follows predictable scaling laws[14].

The authors identify that while applying continuous optimization pressure against a proxy reward model initially improves the true objective (human preference), continued optimization past a critical threshold leads to a severe collapse in true performance[14]. This dynamic is a manifestation of Goodhart's Law: once a measure becomes a target, it ceases to be a reliable measure. In attempting to maximize the proxy reward, the language model learns to exploit spurious correlations, specific stylistic tics, and verbosity rather than generating genuinely accurate or helpful reasoning[14]. To mitigate this, entropy-regularized methods are increasingly employed in reinforcement learning to explicitly reward policy stochasticity, thereby preventing the agent from collapsing onto a single, overoptimized strategy[14].

This overoptimization feeds directly into the "alignment tax"—a documented degradation in a model's generalized zero-shot or few-shot capabilities across standard benchmarks following rigorous safety conditioning[2]. As commercial models are continually updated to patch security vulnerabilities and enforce refusal protocols, the gradients responsible for maintaining complex problem-solving capabilities are penalized, causing the longitudinal performance drift observed by end-users[2].

Main Thematic and Technical Discussion

The theoretical underpinnings of catastrophic forgetting and the alignment tax manifest profoundly when analyzing the empirical behavior of commercial APIs. The degradation is not uniform; instead, it presents as highly volatile, task-specific fluctuations that complicate the deployment of LLMs in production environments.

Empirical Observations of Temporal Drift in Commercial APIs

The most comprehensive investigation into the longitudinal stability of commercial LLMs is documented in the study *How Is ChatGPT's Behavior Changing over Time?* by Chen, Zaharia, and Zou[1]. The researchers systematically evaluated the March 2023 and June 2023 iterations of GPT-3.5 and GPT-4 across diverse computational and linguistic tasks, including mathematical problem-solving, sensitive question answering, code generation, and visual reasoning[1].

The analysis revealed that the behavior of the same LLM service can alter drastically over a span of merely three months, highlighting a critical lack of referential stability[4]. The observed fluctuations underscore the necessity for continuous monitoring, as capabilities both degraded and improved unpredictably across different vectors.

Task Domain	Model	March 2023 Evaluation	June 2023 Evaluation	Documented Behavioral Shift
Prime Number Identification	GPT-4	97.6% Accuracy	2.4% Accuracy	Catastrophic degradation; complete loss of chain-of-thought adherence[4].

Prime Number Identification	GPT-3.5	Poor Accuracy	Substantially Improved	Unpredictable capability gain[4].
Sensitive / Dangerous Questions	GPT-4	High compliance	Drastically reduced compliance	Enhanced safety alignment; significant increase in refusal rates[4].
Code Generation	GPT-4 & GPT-3.5	Clean, executable formatting	High formatting error rate	Increased tendency to output non-executable conversational text surrounding code blocks[4].
Multi-hop Knowledge Questions	GPT-4	Baseline	Improved	Enhanced retrieval and connection of disparate facts[5].
Multi-hop Knowledge Questions	GPT-3.5	Baseline	Degraded	Loss of complex associative memory[5].

The severe degradation in GPT-4's ability to identify prime numbers was primarily attributed to the model's sudden failure to adhere to "chain-of-thought" (CoT) prompting[5]. In the March iteration, the model naturally broke down the prime identification process into sequential logical steps, confirming divisors before generating a conclusion. By June, the model attempted to answer the prompt directly without intermediate reasoning, leading to catastrophic hallucination rates and a failure to discern prime from composite numbers[5].

The formatting mistakes in code generation were similarly driven by a behavioral shift rather than a loss of syntactic knowledge. The newer versions became excessively conversational, wrapping code blocks in extraneous markdown and conversational pleasantries that subsequently broke automated compilation and execution pipelines[4]. However, the academic discourse surrounding this specific dataset suggests nuance. Critics of the study, including Narayanan and Kapoor, argue that the June 2013 (0613) API update specifically enhanced the model's attention to system messages and introduced function-calling capabilities[22]. Therefore, evaluating the June model using standard user prompts optimized for the March model conflates a behavioral shift (expecting instructions in the system message) with a genuine loss of underlying capability[22]. Regardless of the root cause, the operational result for end-users relying on static prompt architectures is severe performance drift.

Trajectories of Readability and Semantic Reliability

The instability of LLM outputs extends beyond discrete benchmarking into the qualitative characteristics of generated text. A longitudinal analysis of authentic texts generated across finance, technology, and public policy domains evaluated outputs using COMET scores, readability indices, sentence length, lexical diversity, and connective incidence[4]. The results revealed no consistent improvement across updates. Instead, semantic reliability and textual characteristics followed diverging trajectories depending on the domain[4]. For instance, while the COMET scores for public policy texts improved over time, the later outputs became significantly less readable, indicating a trade-off between semantic density and human accessibility[4].

Furthermore, performance drift highly impacts specialized evaluations. In comparative diagnostics of medical reasoning capabilities, different state-of-the-art models exhibit wide variance. An evaluation demonstrated that ChatGPT-5 achieved moderate agreement with expert ground truth (accuracy 0.67; kappa = 0.33), whereas Gemini 2.5 Pro demonstrated significantly lower performance and higher uncertainty (accuracy 0.56; kappa = 0.11)[4]. These disparate trajectories emphasize that updates tuned to improve general chat helpfulness often erode highly specialized technical capabilities.

The Illusion of Reasoning and Pattern Matching Susceptibility

A central debate concerning performance drift is whether LLMs are losing actual reasoning capabilities or simply experiencing shifts in their probabilistic pattern-matching heuristics. This question is rigorously explored in *GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models* by Mirzadeh et al.[23].

Historically, mathematical reasoning in LLMs has been assessed using the GSM8K (Grade School Math 8K) dataset. As frontier models approached human parity on this benchmark, researchers questioned whether models were genuinely reasoning or merely memorizing the static training data[23]. To test this, Mirzadeh et al. introduced GSM-Symbolic, generating 100 symbolic templates derived from GSM8K and creating 50 variants per template by dynamically altering proper names, numerical values, and the number of logical clauses[24].

Evaluating 25 state-of-the-art open and closed models using standard eight-shot chain-of-thought prompting revealed massive performance variances that expose the profound fragility of LLM "reasoning"[24].

1. **Numerical and Entity Sensitivity:** Models exhibited high variance simply when the numerical values in a word problem were changed. While altering proper names caused slight perturbations, changing numbers led to significant accuracy drops, suggesting strong token bias and a reliance on learned sequences rather than formal logical execution[23].
2. **Complexity Degradation:** As the number of clauses in a mathematical problem

increased, performance deteriorated at a rate much steeper than the linear increase in logical steps would predict, indicating an inability to maintain long context logic[23].

3. **The GSM-NoOp Vulnerability:** The most compelling evidence against genuine reasoning was the introduction of the GSM-NoOp dataset. By injecting a single, seemingly relevant but mathematically useless clause into the problem (a clause that did not alter the mathematical operations required to reach the answer), state-of-the-art models suffered catastrophic performance drops of up to 65%[23].

The models repeatedly incorporated the irrelevant numbers into their calculations, attempting to force all provided data into a learned heuristic pattern[23]. Providing the models with multiple few-shot examples containing similar distractors did not mitigate the failure[23]. This definitively confirms that LLMs operate via sophisticated probabilistic pattern matching rather than formal logic; they lack the architectural capacity to discern relevant from irrelevant information if both mimic the semantic structure of their training distribution[23]. Consequently, any drift in the underlying training distribution—such as appending new alignment data—drastically alters the model's pattern-matching landscape, causing unpredictable downstream failures in tasks completely unrelated to the update.

Current Approaches and Developments in Evaluation

Because performance drift continuously invalidates static benchmark scores, the methodologies utilized to evaluate LLMs have undergone a profound evolution. The field has shifted from relying on fixed datasets to dynamic, continuously updating ecosystems and automated, AI-driven judiciaries.

Dynamic and Contamination-Free Benchmarking: LiveBench

The phenomenon of test set contamination—where static benchmark questions inadvertently leak into the massive pre-training or fine-tuning datasets of newer models—artificially inflates performance metrics and masks true capability drift[29]. For example, empirical studies demonstrated that LLM performance on Codeforces dropped significantly after the model's training cutoff date, proving that earlier high scores were the result of data memorization rather than generalized coding ability[31]. Similarly, models have been shown to overfit to hand-crafted variants of the established GSM8K dataset[31].

To resolve this, a consortium of researchers released *LiveBench: A Challenging, Contamination-Free LLM Benchmark*[32]. LiveBench operates on a fundamentally different paradigm designed to track model capabilities objectively over time without the risk of data leakage.

The architecture of LiveBench relies on three core operational pillars:

1. **Continuous Content Refresh:** Tasks are sourced from frequently updated and highly recent information streams. Questions are procedurally generated from newly published mathematical competitions, recent arXiv preprints, contemporary news articles, and fresh Kaggle datasets[29]. The benchmark completely refreshes its dataset on a regular 6-month cycle, with monthly delayed releases to prevent scraping[29].
2. **Objective Ground-Truth Scoring:** In stark contrast to subjective benchmarks, every LiveBench question has a verifiable, objective ground-truth answer. This allows complex

problems to be scored automatically via deterministic code execution or exact string matching, completely eliminating the biases inherent in crowdsourced human judging or LLM-based evaluation[29].

3. **Diverse Categorization:** The benchmark currently contains 18 diverse tasks spanning six primary domains: Mathematics, Reasoning, Data Analysis, Language Comprehension, Coding, and Instruction Following[32].

LiveBench has demonstrated that frontier models are significantly less capable than static benchmarks imply. Evaluation results from 1000 samples gathered in August 2024 revealed that top-performing architectures, such as GPT-4o, maintain an average accuracy around 65%, rarely exceeding a 70% threshold across the suite[30].

To extract actionable insights from LiveBench, researchers have developed advanced checklist-driven evaluation and self-improvement protocols, namely TICK and STICK[34]. The TICK protocol uses LLMs to generate and apply YES/NO checklists for each instruction, increasing the agreement between LLM and human judgment from 46.4% to 52.2%[34]. The STICK protocol applies checklist-driven self-refinement; when applied to the Command-R+ model, its reasoning score increased from 29.2 to 37.0 (a +7.8% gain) on LiveBench tasks, demonstrating that highly structured, checklist-based feedback is substantially more effective at overcoming drift than unstructured self-critique[34]. LiveBench scores correlate strongly (0.88–0.91) with other prominent leaderboards like Chatbot Arena, making it a highly reliable metric for evaluating model stability[34]. Similarly, platforms like WildBench benchmark LLMs using challenging, contamination-free tasks sourced directly from real users in the wild, ensuring ecological validity[35].

LLM-as-a-Judge and Evaluator Paradigms

While objective benchmarks like LiveBench are optimal for mathematics and coding, evaluating open-ended chat assistants for alignment, tone, safety, and helpfulness requires nuanced preference ranking. Given the prohibitive cost and slow speed of human evaluation, the industry has widely adopted the "LLM-as-a-Judge" paradigm to track subjective performance drift.

The foundational study establishing this framework is *Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena* by Zheng, Chiang, et al.[15]. The researchers introduced two primary mechanisms to assess open-ended conversational drift: MT-Bench (a multi-turn question set designed to test conversational flow and instruction following) and Chatbot Arena (a crowdsourced, blind, pairwise battle platform for LLMs)[15].

By utilizing a strong frontier model (e.g., GPT-4) as an automated judge to approximate human preferences, Zheng et al. demonstrated an agreement rate exceeding 80% between the LLM judge and human evaluators—a rate statistically equivalent to the inter-annotator agreement found between different human experts[15]. The LLM-as-a-Judge framework offers unparalleled scalability, explainability, and speed, enabling organizations to run continuous regression testing on every pull request, prompt change, or model swap to catch performance drift before production deployment[18].

However, LLM judges possess inherent cognitive biases that must be rigorously mitigated through algorithmic guardrails:

- **Position Bias:** When conducting pairwise evaluations (Model A vs. Model B), LLM judges exhibit a measurable tendency to favor the response presented first, regardless of quality[18]. Mitigation requires running every comparison twice with the response order swapped, and aggregating the results to ensure positional neutrality[39].
- **Verbosity Bias:** Judge models overwhelmingly favor longer, more verbose answers, conflating length with quality[15]. This must be mitigated by explicitly embedding strict conciseness requirements into the evaluation rubric[39].
- **Self-Enhancement Bias:** An LLM judge will disproportionately favor responses generated by its own model architecture (e.g., Claude scoring Claude higher than GPT)[18]. To mitigate this, practitioners employ "ensemble judging"—using multiple, diverse LLMs to evaluate a single response—and down-weight the generator model's score in the final calculation[40].
- **Limited Reasoning Capacity and Calibration:** LLM judges fail profoundly when asked to evaluate complex mathematical or coding questions without a reference. Without ground truth, judges only agree with human experts on questions the judge model itself is capable of solving[37]. Consequently, "reference-guided grading" and prompting the judge to explicitly output its reasoning trace before assigning a final score are mandatory protocols[18]. Furthermore, judges must be continuously calibrated against human annotations using metrics like Cohen's Kappa to ensure they agree with humans beyond random chance[40].

Research Challenges and Limitations

Despite rapid advancements in dynamic evaluation, the ongoing effort to track and mitigate performance drift is constrained by several profound structural limitations within the AI ecosystem.

Exception Chain Collapse and Regulatory Compliance

The shifting baseline of empirical boundaries creates an intractable problem for compliance in regulated industries. In the legal, medical, and financial sectors, an AI application must pass rigorous validation testing. However, performance drift causes what researchers term "exception chain collapse." In evaluations using the Aethis Eligibility Module on nested conditional rules (e.g., "A is required UNLESS B applies, UNLESS C overrides B"), empirical boundaries proved entirely unstable[8]. A frontier model that scored 96.6% on a construction insurance logic task silently shifted to 100% accuracy weeks later without a version bump, while concurrently failing newly introduced edge cases it had previously navigated successfully[8]. If an API provider updates the model and causes an exception chain collapse, a previously certified workflow is instantly rendered non-compliant.

Latent Space Opacity

The opacity of closed-source commercial APIs prevents proper root-cause analysis of drift. Because model providers silently update weights, introduce latent quantization for inference speed, and alter underlying Mixture of Experts routing systems without transparency, independent researchers are left treating the model as a black box[1]. It is often impossible to decouple whether an observed performance drop is due to parameter-level catastrophic forgetting, a new hidden system prompt overriding user instructions, or hardware-level optimization[22].

Knowledge Conflict and Confident Conflicts

A core mathematical challenge causing drift during alignment is the phenomenon of "knowledge conflict." During Supervised Fine-Tuning, models frequently encounter "Confident Conflicts"—instances where the model's pre-trained priors express high confidence (low entropy) in a specific output, but the new fine-tuning data forces it to learn a divergent ground truth[2]. Standard fine-tuning protocols force the model to fit these conflicts based purely on probability, triggering highly destructive gradient updates that overwrite broad swaths of foundational knowledge[2]. While experimental approaches like Entropy-Adaptive Fine-Tuning (EAFT) attempt to use token-level entropy as a gating mechanism to suppress gradients on conflicting data, scaling these interventions to massive, distributed architectures remains a profound computational hurdle[2].

Research Gaps and Future Directions

The current literature suggests several critical avenues for future research to stabilize LLM performance and build resilient artificial intelligence systems capable of continuous learning without continuous forgetting.

Referential Security and AI Identity

There is an urgent need to establish "referential security" as a core paradigm in AI evaluation[4]. Currently, safety and capability claims are ephemeral; a benchmark claim made about a commercial model in March may be entirely invalid by June. Future research must develop cryptographic or architectural mechanisms to lock, version, and verify the precise identity and state of a model at inference time. This approach reframes model identity as an empirically verifiable property, ensuring that downstream auditing parties can conclusively determine which system a specific safety claim addressed, insulating workflows from silent drift[4].

Neuro-Symbolic Integration and Illocutionary Planning

To bypass the inherent fragility of autoregressive reasoning demonstrated in GSM-Symbolic, future architectures must decouple deterministic execution from probabilistic text generation[8]. Integrating LLMs with Symbolic Math Solvers (SMT engines) or deterministic code executors ensures that once a rule is authored by the LLM, its execution is mathematically guaranteed, insulating the output from the semantic drift of the LLM's latent space[8]. The Aethis Eligibility Module demonstrated this by pairing an LLM authoring layer with a deterministic execution engine, scoring a perfect 20/20 on an adversarial construction insurance task where state-of-the-art standalone LLMs failed entirely due to silent updates[8]. Validation on LegalBench across 949 held-out cases further proved the superiority of this decoupled approach[8]. Similarly, to combat source hallucination and drift, researchers have introduced "chain-of-illocution" (CoI) prompting, motivated by Achinstein's illocutionary theory of explanation. CoI expands queries into implicit explanatory questions that drive retrieval, yielding statistically significant gains of up to 63% in source adherence[4].

LLM-SLM Collaboration and Federated Architectures

To address the latency, privacy, and drift issues of massive cloud-based LLMs, researchers are pivoting toward collaborative architectures combining Large Language Models with Small Language Models (SLMs)[41]. In this paradigm, SLMs handle edge-device routing, speculative decoding, and local adaptation, while the LLM handles complex generalization[42]. Managing drift in these distributed environments requires novel algorithmic approaches, such as the FedDAT framework. FedDAT facilitates federated instruction tuning across decentralized private data while utilizing domain coverage augmentation and adapter mechanisms to balance global knowledge and mitigate client-specific model drift[43]. Furthermore, frameworks like FrugalGPT allow for cascaded queries across multiple models, significantly reducing operational costs while dynamically routing around the performance drift of any single API[5].

Economic World Models and Continuous Drift Monitoring

The expansion of AI into complex financial and economic simulations requires the development of Economic World Models (EWMs) and platforms like FinGPT, which rely on the continuous, lightweight adaptation of financial models[45]. In these high-stakes multi-turn environments, static snapshot guardrails are insufficient. Future evaluation platforms will employ continuous survival analysis to map the latent space of models over time. The Triple-tier Anomaly Defense (TRIAD) framework models safety verification as a dynamic survival prediction problem[19]. By treating multi-turn conversational flow as a continuous topological trajectory, TRIAD utilizes a Bayesian Hidden Markov Model feedback loop and a regularized Mahalanobis distance metric to monitor covariance shifts in high-dimensional spaces[19]. This enables the system to differentiate between benign conversational exploration and malicious structural poisoning, calculating an expected time-to-failure to halt catastrophic semantic drift before policy violations occur[19].

Conclusion

The temporal instability of Large Language Models represents one of the most critical systemic barriers to their safe integration into high-stakes scientific, commercial, and regulatory environments. As evidenced by exhaustive longitudinal studies, performance drift is not a rare anomaly but an inherent feature of continuous post-training. This instability is driven by the mathematical realities of catastrophic forgetting during sequential fine-tuning, the alignment tax incurred by safety conditioning, and the scaling laws of reward model overoptimization. The revelation that current state-of-the-art models rely heavily on probabilistic pattern matching rather than formal logic—demonstrated by their catastrophic failure upon the introduction of irrelevant variables in symbolic mathematical benchmarks—underscores their profound vulnerability to seemingly benign algorithmic updates.

To counteract this, the paradigm of AI evaluation is undergoing a fundamental transformation. Static, easily memorized benchmarks are being systematically replaced by contamination-resistant, continually updating frameworks like LiveBench, which mandate objective ground-truth scoring and procedural question generation. Simultaneously, the deployment of carefully calibrated LLM-as-a-Judge pipelines provides the scalable infrastructure necessary to continuously monitor subjective conversational drift at the speed of deployment. Ultimately, stabilizing the performance of artificial intelligence will require moving beyond pure autoregressive scaling toward neuro-symbolic architectures, referential security paradigms, and targeted algorithmic interventions like entropy-adaptive fine-tuning that dynamically shield foundational knowledge from the destructive forces of continuous alignment.

Complete References

[1] Chen, L., Zaharia, M., & Zou, J. (2023). How Is ChatGPT's Behavior Changing over Time? arXiv:2307.09009.

[6] Longitudinal study evaluating GPT-4o performance on physics tasks over time. arXiv:2602.15889.

[29] White, C., Dooley, S., Goldblum, M., et al. (2024). LiveBench: A Challenging, Contamination-Free LLM Benchmark. International Conference on Machine Learning. arXiv:2406.19314.

[34] Shi, et al. (2025). InferenceDynamics.

[35] Lin, B. Y., Deng, Y., Chandu, K., Brahman, F., Ravichander, A., Pyatkin, V., Dziri, N., Le Bras, R., & Choi, Y. (2024). WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild. arXiv:2406.04770.

[51] Nagireddy, M., Sattigeri, P., et al. (2024/2025). Model drift and IBM AI Ethics detectors. Association for the Advancement of Artificial Intelligence (AAAI). arXiv:2403.06009.

[45] Research on Economic World Models (EWMs) and FinGPT. arXiv:2608.06020.

[2] Research on Entropy-Adaptive Fine-Tuning (EAFT) to mitigate catastrophic forgetting. arXiv:2601.02151.

[8] Research detailing the Aethis Eligibility Module and exception chain collapse. arXiv:2607.23386.

[41] Survey on Large Language Model and Small Language Model (LLM-SLM) Collaboration. arXiv:2505.07460.

[19] Research on the Triple-tier Anomaly Defense (TRIAD) framework. arXiv:2605.18988.

[43] FedDAT: Federated instruction tuning of LLMs. arXiv:2503.16585.

[3] Luo, Y., Yang, Z., Meng, F., Li, Y., Zhou, J., & Zhang, Y. (2023/2025). An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning. IEEE Transactions on Audio Speech and Language Processing, 33, 3776-3786. arXiv:2308.08747.

[14] Gao, L., Schulman, J., & Hilton, J. (2023). Scaling Laws for Reward Model Overoptimization. Proceedings of the 40th International Conference on Machine Learning (ICML'23). arXiv:2112.09332.

[23] Mirzadeh, I., Alizadeh, K., Shahrokhi, H., Tuzel, O., Bengio, S., & Farajtabar, M. (2024/2025). GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models. ICLR 2025. arXiv:2410.05229.

[53] Zhou, J., Lu, T., Hou, L., et al. (2023). Instruction-Following Evaluation for Large Language Models. arXiv:2311.07911.

[15] Zheng, L., Chiang, W.-L., Li, Z., Lin, Z., Sheng, Y., Wu, Z., Zhang, H., Zhuang, S., Zhuang, Y., Gonzalez, J. E., Stoica, I., & Xing, E. P. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. Neural Information Processing Systems (NeurIPS 2023). arXiv:2306.05685.

Works cited

1. ChatGPT Behavior Changes Over Time | PDF | Prime Number - Scribd, https://www.scribd.com/document/659959844/How-is-ChatGPT-s-Behaviour-Changing-Over-Time-2307-09009
2. Entropy-Adaptive Fine-Tuning: Resolving Confident Conflicts to Mitigate Forgetting - arXiv, https://arxiv.org/html/2601.02151v1
3. An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-Tuning | Request PDF - ResearchGate, https://www.researchgate.net/publication/395308608_An_Empirical_Study_of_Catastrophic_Forgetting_in_Large_Language_Models_During_Continual_Fine-Tuning
4. How is ChatGPT's behavior changing over time? - ResearchGate, https://www.researchgate.net/publication/372445224_How_is_ChatGPT's_behavior_changing_over_time
5. Lingjiao Chen - alphaXiv, https://www.alphaxiv.org/@lingjiao-chen
6. Daily and Weekly Periodicity in Large Language Model Performance and Its Implications for Research - arXiv, https://arxiv.org/html/2602.15889v2
7. [2307.09009] How is ChatGPT's behavior changing over time? - arXiv, https://arxiv.org/abs/2307.09009
8. Confidently Wrong: Exception Chain Collapse in Frontier LLM Rule Evaluation - arXiv, https://arxiv.org/html/2607.23386v1
9. An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning | alphaXiv, https://www.alphaxiv.org/abs/2308.08747
10. Mechanistic Analysis of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning - arXiv, https://arxiv.org/html/2601.18699v1
11. [2308.08747] An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning - arXiv, https://arxiv.org/abs/2308.08747
12. [2308.08747] An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-tuning - ar5iv, https://ar5iv.labs.arxiv.org/html/2308.08747
13. An Empirical Study of Catastrophic Forgetting in Large Language Models During Continual Fine-Tuning - Semantic Scholar, https://www.semanticscholar.org/paper/An-Empirical-Study-of-Catastrophic-Forgetting-in-Luo-Yang/838cd69a0b6c9c244a6eebb0f4742c0625132de6
14. Scaling Laws for Reward Model Overoptimization - ResearchGate, https://www.researchgate.net/publication/364436070_Scaling_Laws_for_Reward_Model_Overoptimization
15. (PDF) Judging LLM-as-a-judge with MT-Bench and Chatbot Arena - ResearchGate, https://www.researchgate.net/publication/371490428_Judging_LLM-as-a-judge_with_MT-Bench_and_Chatbot_Arena
16. Reward Model Overoptimisation in Iterated RLHF - arXiv, https://arxiv.org/html/2505.18126v1
17. Jacob Hilton - Google Scholar, https://scholar.google.co.in/citations?user=WyKvz7EAAAAJ&hl=vi

18. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena - arXiv, https://arxiv.org/html/2306.05685v4
19. Surviving the Unseen: Predictive Defense for Novel Multi-Turn Multimodal Attacks - arXiv, https://arxiv.org/pdf/2605.18988
20. How Is ChatGPT's Behavior Changing over Time? - arXiv, https://arxiv.org/pdf/2307.09009
21. How Is ChatGPT's Behavior Changing over Time? | by Farzad Karami - Medium, https://medium.com/@farzad.karami/how-is-chatgpts-behavior-changing-over-time-3f2d67b7d045
22. Paper page - How is ChatGPT's behavior changing over time? - Hugging Face, https://huggingface.co/papers/2307.09009
23. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models - arXiv, https://arxiv.org/pdf/2410.05229
24. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models — Read with an AI tutor - ChapterPal, https://www.chapterpal.com/s/vv6paebr/gsm-symbolic-understanding-the-limitations-of-mathematical-reasoning-in-large-language-models
25. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models | OpenReview, https://openreview.net/forum?id=AjXkRZIvjB
26. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models | Request PDF - ResearchGate, https://www.researchgate.net/publication/384698709_GSM-Symbolic_Understanding_the_Limitations_of_Mathematical_Reasoning_in_Large_Language_Models
27. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models - arXiv, https://arxiv.org/html/2410.05229v1?utm_source
28. GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models - arXiv, https://arxiv.org/html/2410.05229v2
29. LiveBench: A Challenging, Contamination-Free LLM Benchmark - GitHub, https://github.com/livebench/livebench
30. LiveBench: A Challenging, Contamination-Limited LLM Benchmark - Semantic Scholar, https://www.semanticscholar.org/paper/LiveBench%3A-A-Challenging%2C-Contamination-Limited-LLM-White-Dooley/774d01e152003f342596031c0c0fbf1936dee41a
31. livebench:achallenging, contamination-limited llm benchmark - arXiv, https://arxiv.org/pdf/2406.19314
32. LiveBench: A Challenging, Contamination-Free LLM Benchmark – Inspect Evals, https://ukgovernmentbeis.github.io/inspect_evals/evals/knowledge/livebench/
33. LiveBench: A Challenging, Contamination-Limited LLM Benchmark | OpenReview, https://openreview.net/forum?id=sKYHBTAxVa
34. LiveBench: Dynamic LLM Benchmark Suite - Emergent Mind, https://www.emergentmind.com/topics/livebench
35. allenai/WildBench: Benchmarking LLMs with Challenging Tasks from Real Users - GitHub, https://github.com/allenai/WildBench

36. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena | OpenReview, https://openreview.net/forum?id=uccHPGDlao
37. Judging LLM-as-a-judge with MT-Bench and Chatbot Arena - Semantic Scholar, https://www.semanticscholar.org/paper/Judging-LLM-as-a-judge-with-MT-Bench-and-Chatbot-Zheng-Chiang/a0a79dad89857a96f8f71b14238e5237cbfc4787
38. Commit-Aware Learning-Based Test Case Prioritization for Continuous Integration - arXiv, https://arxiv.org/html/2604.25363v1
39. LLM-as-a-Judge: How to Build Reliable, Scalable Evaluation for LLM Apps and Agents, https://www.comet.com/site/blog/llm-as-a-judge/
40. How to Use LLM as a Judge (Without Getting Burned) - Manthan Gupta, https://manthanguptaa.in/posts/llm_as_a_judge/
41. A Survey on Collaborative Mechanisms Between Large and Small Language Models - arXiv, https://arxiv.org/html/2505.07460v1
42. A Survey on Collaborative Mechanisms Between Large and Small Language Models - arXiv, https://arxiv.org/pdf/2505.07460
43. Distributed LLMs and Multimodal Large Language Models: A Survey on Advances, Challenges, and Future Directions - arXiv, https://arxiv.org/pdf/2503.16585
44. Distributed LLMs and Multimodal Large Language Models: A Survey on Advances, Challenges, and Future Directions - arXiv, https://arxiv.org/html/2503.16585v1
45. From Economic Agents to Agentic Economies: A Systems Blueprint for Economic World Models - arXiv, https://arxiv.org/html/2608.06020v1
46. How is ChatGPT's behavior changing over time[Paper Explained] - YouTube, https://www.youtube.com/watch?v=7vsIA4_jIY0
47. Lingjiao Chen's Homepage, http://lingjiaochen.com/publications.html
48. How is ChatGPT's behavior changing over time? - dblp, https://dblp.org/rec/journals/corr/abs-2307-09009.html
49. LiveBench Leaderboard - LLM Stats, https://llm-stats.com/benchmarks/livebench
50. [2406.19314] LiveBench: A Challenging, Contamination-Limited LLM Benchmark - arXiv, https://arxiv.org/abs/2406.19314
51. arXiv:2403.06009v3 [cs.LG] 19 Aug 2024, https://arxiv.org/pdf/2403.06009
52. [2410.05229] GSM-Symbolic: Understanding the Limitations of Mathematical Reasoning in Large Language Models - arXiv, https://arxiv.org/abs/2410.05229
53. [PDF] Instruction-Following Evaluation for Large Language Models - Semantic Scholar, https://www.semanticscholar.org/paper/Instruction-Following-Evaluation-for-Large-Language-Zhou-Lu/1a9b8c545ba9a6779f202e04639c2d67e6d34f63
54. LexInstructEval: Lexical Instruction Following Evaluation for Large Language Models - arXiv, https://arxiv.org/html/2511.17561v2
55. Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena, https://www.chapterpal.com/s/3ha9zzte/judging-llm-as-a-judge-with-mt-bench-and-chatbot-arena